

Designing a Neural Network

Joe Tilley

August 14, 2017

Abstract

This paper is based on Michael Nielsen's online book, "Neural Networks and Deep Learning" [1]. I will synthesise Nielsen's book, approaching from a more mathematical angle and using original codes and examples.

1 An Introduction to a basic Neural Network

The aim of a neural network is to act as a complicated function, taking many inputs and producing a variety of outputs. Nielsen uses the example of a neural network understanding handwritten digits, taking inputs of the darkness of each pixel in the picture of the digit and producing an output of the probability of it being each digit.

1.1 Setting up the Network

The network is made up of perceptrons. A *perceptron* a function of the form

$$p(x_1, x_2, \dots, x_n) = f\left(\sum_{i=1}^n w_i x_i + b\right)$$

where w_i, b are the *weights* and *bias*, chosen by the user. For notational purposes, we let \mathbf{x}, \mathbf{w} be the column vectors of inputs and weights respectively. Hence,

$$p(\mathbf{x}) = f(\mathbf{w} \cdot \mathbf{x} + b)$$

The function f , called the *activation function*, can be chosen, but in this section, we will let f be the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

We create a network of perceptron layers, each layer taking every input from the previous layer and acting as an input to the next (i.e. given two layers next to each other, we form a complete bipartite graph). If the function f is continuous, then a small change in any single w_i or b will result in a small change in the output, since the network is a composition of continuous functions.

We can encode this system in terms of matrices. Suppose we are given layer ℓ with n perceptrons and layer $\ell + 1$ with m perceptrons. The n perceptrons will have an associated output x_1, \dots, x_n , called the *activation output from layer ℓ* , and each of the n perceptrons will have m associated weights, perceptron $1 \leq i \leq n$ in level ℓ having weights w_{1i}, \dots, w_{mi} . Then, the column vector of activation outputs from level $\ell + 1$ is given by

$$\mathbf{x}_{\ell+1} = \sigma(\mathbf{w}_{\ell+1} \cdot \mathbf{x}_{\ell} + \mathbf{b}_{\ell+1})$$

where taking σ of a column vector is interpreted as taking σ of each of the values. Here, \mathbf{x}_{ℓ} is the column vector of n activation outputs from level ℓ , $\mathbf{b}_{\ell+1}$ is the column vector of biases for level $\ell + 1$, and the weights

$\mathbf{w}_{\ell+1} = (w_{ij}^{(\ell+1)})$ is an $m \times n$ matrix where $w_{ij}^{(\ell+1)}$ is the weight between perceptron j in level ℓ to perceptron i in level $\ell + 1$.

Based on Nielsen's code, I have written a class which builds the network with random, normally distributed weights and biases. Moreover, the function feedforward is shown, which will take the initial input vector \mathbf{x}_1 and compute the activation output vector \mathbf{x}_ℓ (where ℓ is the number of layers) given predetermined weights and biases. We compute this value through, what appears to be, quite a contrived value called z , which we also store and output by the function. Its importance will only be shown later in the section 'Backpropagation', so feel free to ignore it for now.

```
class Network(object):

    def __init__(self, sizes):
        self.sizes = sizes
        self.L = len(sizes)
        self.biases = [np.random.normal(0, 1, (sizes[n+1], 1)) for n in range(len(sizes)-1)]
        self.weights = [np.random.normal(0, 1, (sizes[n+1], sizes[n])) for n in range(len(sizes)-1)]

    def feedforward(self, input_vector):
        ''' Runs the input vector through the network'''
        z = [np.add(self.weights[0].dot(input_vector), self.biases[0])]
        for n in range(1, self.L-1):
            z.append(np.add(self.weights[n].dot(sigmoid(z[-1])), self.biases[n]))
        output_vector = sigmoid(z[-1])
        return [z, output_vector]

def sigmoid(x):
    return 1/(1+np.exp(-x))
```

1.2 Training the Network

The aim of our neural network is for it to learn from its previous attempts. We shall have some training data, T , from which it shall learn. T will contain tuples $(\mathbf{x}, \mathbf{y}_\mathbf{x})$ where \mathbf{x} is the input vector and $\mathbf{y}_\mathbf{x}$ the true output vector associated to \mathbf{x} (which we wish our network to produce). Our aim is to minimise the function

$$C(\mathbf{w}, \mathbf{b}) = \frac{1}{2n} \sum_{\mathbf{x} \in T} \|N(\mathbf{w}, \mathbf{b}, \mathbf{x}) - \mathbf{y}_\mathbf{x}\|^2$$

where $n = |T|$ is the number of training tuples, and $N(\mathbf{w}, \mathbf{b}, \mathbf{x})$ is the output of the network given those weights, biases and input vector, i.e. \mathbf{x}_ℓ the activation output from the final layer ℓ . We minimise this function by our choice of \mathbf{w} and \mathbf{b} . Other cost functions can be chosen, and we will explore this later.

The function is too complicated to find the global derivative of, so we use a method of steepest gradient descent. Suppose we are given weights and biases $(\mathbf{w}_n, \mathbf{b}_n)$, and we want to iterate to find a new set of weights and biases such that

$$C(\mathbf{w}_{n+1}, \mathbf{b}_{n+1}) \leq C(\mathbf{w}_n, \mathbf{b}_n)$$

Gradient descent tell us that if we let

$$(\mathbf{w}_{n+1}, \mathbf{b}_{n+1}) = (\mathbf{w}_n, \mathbf{b}_n) - \eta \nabla C(\mathbf{w}_n, \mathbf{b}_n)$$

for sufficiently small $\eta > 0$, the *learning rate*. This will work, since

$$\Delta C = \nabla C \cdot \mathbf{x} = \nabla C \cdot (-\eta \nabla C) = -\eta \|\nabla C\|^2 \leq 0$$

Our task now is to compute ∇C , the gradient of C . However, this is no easy task, since we will typically have large T . Instead, we will complete this by a method of averaging stochastic gradient descent. We take a small sample $T' \subset T$ (this sample could even be size of just one), to get

$$\nabla C \approx \frac{1}{|T'|} \sum_{\mathbf{x} \in T'} \nabla C_{\mathbf{x}}$$

where $C_{\mathbf{x}} = \frac{1}{2} \|N(\mathbf{w}, \mathbf{b}, \mathbf{x}) - \mathbf{y}_{\mathbf{x}}\|^2$. In particular, the iteration for any given weight w_n from the weights \mathbf{w}_n or bias b_n from the biases \mathbf{b}_n is given by

$$w_{n+1} = w_n - \frac{\eta}{|T'|} \sum_{\mathbf{x} \in T'} \frac{\partial C_{\mathbf{x}}}{\partial w_n}$$

and

$$b_{n+1} = b_n - \frac{\eta}{|T'|} \sum_{\mathbf{x} \in T'} \frac{\partial C_{\mathbf{x}}}{\partial b_n}$$

The question now is: how do we calculate the values of $\frac{\partial C_{\mathbf{x}}}{\partial w_n}$ and $\frac{\partial C_{\mathbf{x}}}{\partial b_n}$? This will introduce an entirely new chapter, so we shall present the code for stochastic gradient descent, leaving blank these values which we will fill in later.

```
def SGD(self, training_data, epochs, mini_batch_size, eta, test_data):
    '''Completes the method of stochastic gradient descent'''
    for n in range(epochs):
        mini_batch = random.sample(training_data, mini_batch_size)
        '''We need to replace (?) with the average (over minibatch) derivative of C
        wrt each term in bias/weight'''
        [dCdb, dCdw] = (?)
        self.biases = [np.add(b, np.multiply(-(eta/mini_batch_size), Cb)) for b, Cb in
                        zip(self.biases, dCdb)]
        self.weights = [np.add(w, np.multiply(-(eta/mini_batch_size), Cw)) for w, Cw
                        in zip(self.weights, dCdw)]
        print("Epoch {0}: {1} / {2}".format(n, self.evaluate(test_data), len(
            test_data)))

def evaluate(self, test_data):
    '''Calculates number of correct answers network produces for the test data'''
    correct = 0
    for (x,y) in test_data:
        network_ans = self.feedforward(x)[1]
        real_ans = y
        if network_ans.index(max(network_ans)) == real_ans.index(max(real_ans)):
            xy_works = 1
        else:
            xy_works = 0
        correct += xy_works
    return correct
```

Note I have also included an evaluate function, which allows you to compare the correct answer against what the neural network has output. The truth condition can be changed depending on the expected output. Nielsen calls an answer ‘correct’ if the highest value of the output vector is in the same index as the ideal output vector, which seems suitable for many contexts, hence I have used this, though it need not always be the case and any condition that the user wants can be chosen depending on the context.

1.3 Backpropagation

In this section, we will find a way to compute the partial derivatives $\frac{\partial C_{\mathbf{x}}}{\partial \mathbf{w}}$ and $\frac{\partial C_{\mathbf{x}}}{\partial \mathbf{b}}$. For notational purposes, we will write $C_{\mathbf{x}}$ simply as C . Recall that given layer ℓ with n perceptrons and layer $\ell+1$ with m perceptrons, the neural network works by the system of equations

$$\mathbf{x}_{\ell} = \sigma(\mathbf{w}_{\ell} \cdot \mathbf{x}_{\ell-1} + \mathbf{b}_{\ell})$$

We define the quantity

$$\mathbf{z}_{\ell} = \mathbf{w}_{\ell} \cdot \mathbf{x}_{\ell-1} + \mathbf{b}_{\ell}$$

and call this the *weight input to layer ℓ* . This was used in the programming of the feedforward algorithm. This definition will play an important role in finding the derivatives required. Furthermore, we define the *error of layer ℓ* by

$$\boldsymbol{\delta}_{\ell} = (\delta_i^{\ell}) \quad \text{where} \quad \delta_i^{\ell} = \frac{\partial C}{\partial z_i^{\ell}}$$

where z_i^{ℓ} is the i^{th} entry of \mathbf{z}_{ℓ} , i.e. the weight input to neuron i of layer ℓ . So,

$$\mathbf{x}_{\ell} = \sigma(\mathbf{z}_{\ell})$$

Firstly, we will find a formula for $\boldsymbol{\delta}_L$ the error of the final layer L , and then use this to compute all other errors, and finally the required partial derivatives. Using multivariate chain rule, we see

$$\delta_i^L = \frac{\partial C}{\partial z_i^L} = \sum_k \frac{\partial C}{\partial x_k^L} \frac{\partial x_k^L}{\partial z_i^L}$$

However, we know that

$$x_k^L = \sigma(z_k^L) \quad \text{so} \quad \frac{\partial x_k^L}{\partial z_i^L} = \begin{cases} \sigma'(z_i^L) & \text{for } i = k \\ 0 & \text{else} \end{cases}$$

Hence

$$\delta_i^L = \frac{\partial C}{\partial x_i^L} \sigma'(z_i^L)$$

We can simplify this further, since

$$C = \frac{1}{2} \|\mathbf{N}(\mathbf{w}, \mathbf{b}, \mathbf{x}) - \mathbf{y}_{\mathbf{x}}\|^2 = \frac{1}{2} \|\mathbf{x}_L - \mathbf{y}_{\mathbf{x}}\|^2 \quad \text{so} \quad \frac{\partial C}{\partial x_k^L} = x_k^L - y_k$$

and in vector form,

$$\boldsymbol{\delta}_L = (\mathbf{x}_L - \mathbf{y}_{\mathbf{x}}) \circ \sigma'(\mathbf{z}_L)$$

where \circ represents entrywise product. We now form an iteration to compute the other errors. Using chain rule again, we see

$$\delta_i^{\ell} = \frac{\partial C}{\partial z_i^{\ell}} = \sum_k \frac{\partial C}{\partial z_k^{\ell+1}} \frac{\partial z_k^{\ell+1}}{\partial z_i^{\ell}} = \sum_k \delta_k^{\ell+1} \frac{\partial z_k^{\ell+1}}{\partial z_i^{\ell}}$$

We can compute this derivative, since, after unfolding the matrix definition of weight input into index terms, we know

$$z_k^{\ell+1} = \sum_j w_{kj}^{\ell+1} x_j^{\ell} + b_k^{\ell+1} = \sum_j w_{kj}^{\ell+1} \sigma(z_j^{\ell}) + b_k^{\ell+1}$$

so taking derivatives, we get

$$\frac{\partial z_k^{\ell+1}}{\partial z_i^{\ell}} = w_{ki}^{\ell+1} \sigma'(z_i^{\ell})$$

And so,

$$\delta_i^\ell = \sum_k \delta_k^{\ell+1} w_{ki}^{\ell+1} \sigma'(z_i^\ell)$$

or in vector form,

$$\delta_\ell = (\mathbf{w}_{\ell+1}^T \cdot \delta_{\ell+1}) \circ \sigma'(\mathbf{z}_\ell)$$

where T represents the transpose. The last layer error and this reverse iteration will allow us to compute error in every layer, and this will allow us to compute the partial derivatives. Again, we use the chain rule and differentiating the index terms for the definition of weight input to get

$$\frac{\partial C}{\partial w_{ij}^\ell} = \sum_k \frac{\partial C}{\partial z_k^\ell} \frac{\partial z_k^\ell}{\partial w_{ij}^\ell} = \frac{\partial C}{\partial z_i^\ell} x_j^{\ell-1} = \delta_i^\ell x_j^{\ell-1}$$

and

$$\frac{\partial C}{\partial b_i^\ell} = \sum_k \frac{\partial C}{\partial z_k^\ell} \frac{\partial z_k^\ell}{\partial b_i^\ell} = \frac{\partial C}{\partial z_i^\ell} = \delta_i^\ell$$

Below, I have programmed the algorithm to compute these derivatives. In the previous section, we simply replace the (?) with this algorithm, and we have now completed our neural network. I have also included a sigmoid prime function (the derivative of sigmoid) and a derivative of the cost function with respect to the output vector. This will allow for ease of code modification if a different sigmoid and cost function were chosen.

```
def backpropogation(self, mini_batch):
    '''Calculates sum of partial derivatives of C wrt weights and biases for a
    given minibatch'''
    dCdb = [np.zeros(b.shape) for b in self.biases]
    dCdw = [np.zeros(w.shape) for w in self.weights]
    for (x,y) in mini_batch:
        [z,x_L] = self.feedforward(x)
        activations = [sigmoid(zed) for zed in z]
        activations.insert(0,x)
        delta = [np.multiply(dCdx_L(x_L, y), sigmoid_prime(z[-1]))]
        for n in range(self.L-2,0,-1):
            delta.insert(0,np.multiply(((self.weights[n]).transpose()).dot(delta
            [0]), sigmoid_prime(z[n-1])))
        db = delta
        dw = [np.zeros(w.shape) for w in self.weights]
        for l in range(self.L-1):
            for i in range(self.sizes[l+1]):
                for j in range(self.sizes[l]):
                    dw[l][i][j] = delta[l][i]*activations[l][j]
        dCdb = [np.add(Cb,b) for Cb, b in zip(dCdb, db)]
        dCdw = [np.add(Cw,w) for Cw, w in zip(dCdw, dw)]
    return [dCdb, dCdw]

def sigmoid_prime(x):
    return np.divide(np.exp(x), ((np.exp(x)+1)**2))

def dCdx_L(x, y):
    ''' Calculates derivatives wrt x_L of C.
    x is column vector of activation outputs, x_L.
    y is column vector of true outputs.'''
    return np.add(x,-y)
```

References

- [1] Michael Nielsen. “Neural Networks and Deep Learning”, 2017. URL <http://neuralnetworksanddeeplearning.com/index.html>. [Accessed: 6th July 2017].